

# IMPLEMENTING NEURAL NETWORKS IN PYTHON

LECTURE 6  
SECTION 3  
JUNE 11



ISTITUTE FOR APPLIED  
COMPUTATIONAL SCIENCE  
AT HARVARD UNIVERSITY



UNIVERSITY of  
RWANDA

## Keras A PYTHON LIBRARY FOR NEURAL NETWORKS:

Keras is a python library that provide intuitive api's for building neural networks very quickly.

### Modules we need for basic neural networks

```
#keras model for feedforward neural networks
from keras.models import Sequential ← for defining neural networks layer by layer.

#keras model for layers in feedforward networks
from keras.layers import Dense ← for connecting the nodes from each layer to
                                every node in the next.

#keras model for optimizing training objectives
from keras import optimizers ← for training neural networks
```

# BUILDING A NEURAL NETWORK FOR REGRESSION:

```
#instantiate a feedforward model
model = Sequential()

#add layers sequentially

#input layer: 2 input dimensions
model.add(Dense(2, input_dim=2, activation='relu',
                kernel_initializer='random_uniform',
                bias_initializer='zeros'))

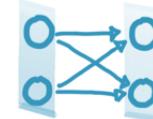
#hidden layer: 2 nodes
model.add(Dense(2, activation='relu',
                kernel_initializer='random_uniform',
                bias_initializer='zeros'))

#output layer: 1 output dimension
model.add(Dense(1, activation='linear',
                kernel_initializer='random_uniform',
                bias_initializer='zeros'))

#configure the model: specify training objective and training algorithm
sgd = optimizers.SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True)
model.compile(optimizer=sgd,
              loss='mean_squared_error')
```

Stochastic gradient descent, learning rate= 0.01  
Loss function for training

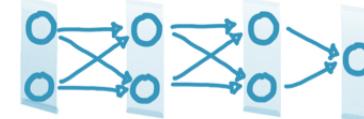
input hidden1



input hidden1 hidden2



input hidden1 hidden2 output



# TRAINING A NEURAL NETWORK:

## Fitting the Model

```
#fit the model and return the mean squared error during training  
history = model.fit(X_train, Y_train, batch_size=20, shuffle=True, epochs=100, verbose=0)
```

→ Compute the gradient using random samples of 20 data points at a time.

→ How many times we should go thru the data set 20 samples at a time

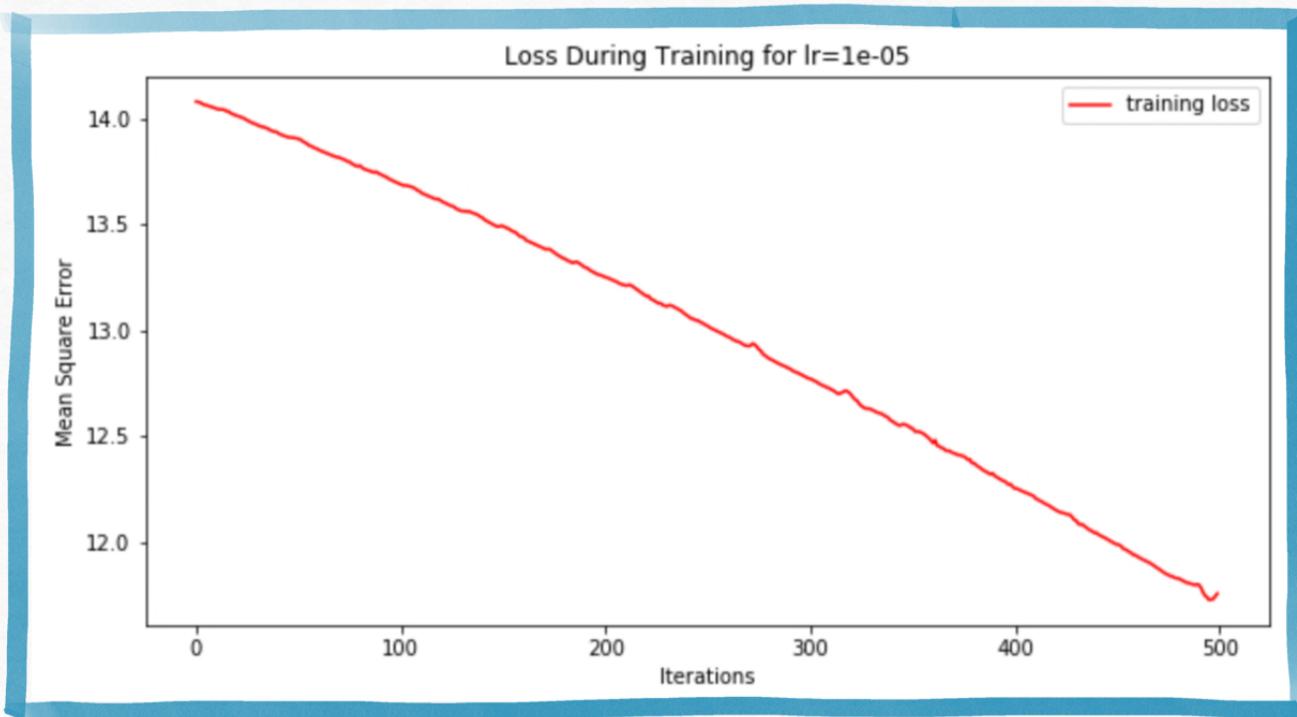
## Visualizing the Trace plot

```
#fit the model and return the mean squared error during training  
history = model.fit(X_train, Y_train, batch_size=20, shuffle=True, epochs=100, verbose=0)  
  
# Plot the loss function and the evaluation metric over the course of training  
fig, ax = plt.subplots(1, 1, figsize=(10, 5))  
  
ax.plot(np.array(history.history['mean_squared_error']), color='blue', label='training accuracy')  
  
plt.show()
```

→ the loss function value is stored under the key 'mean\_squared\_error' in history.history

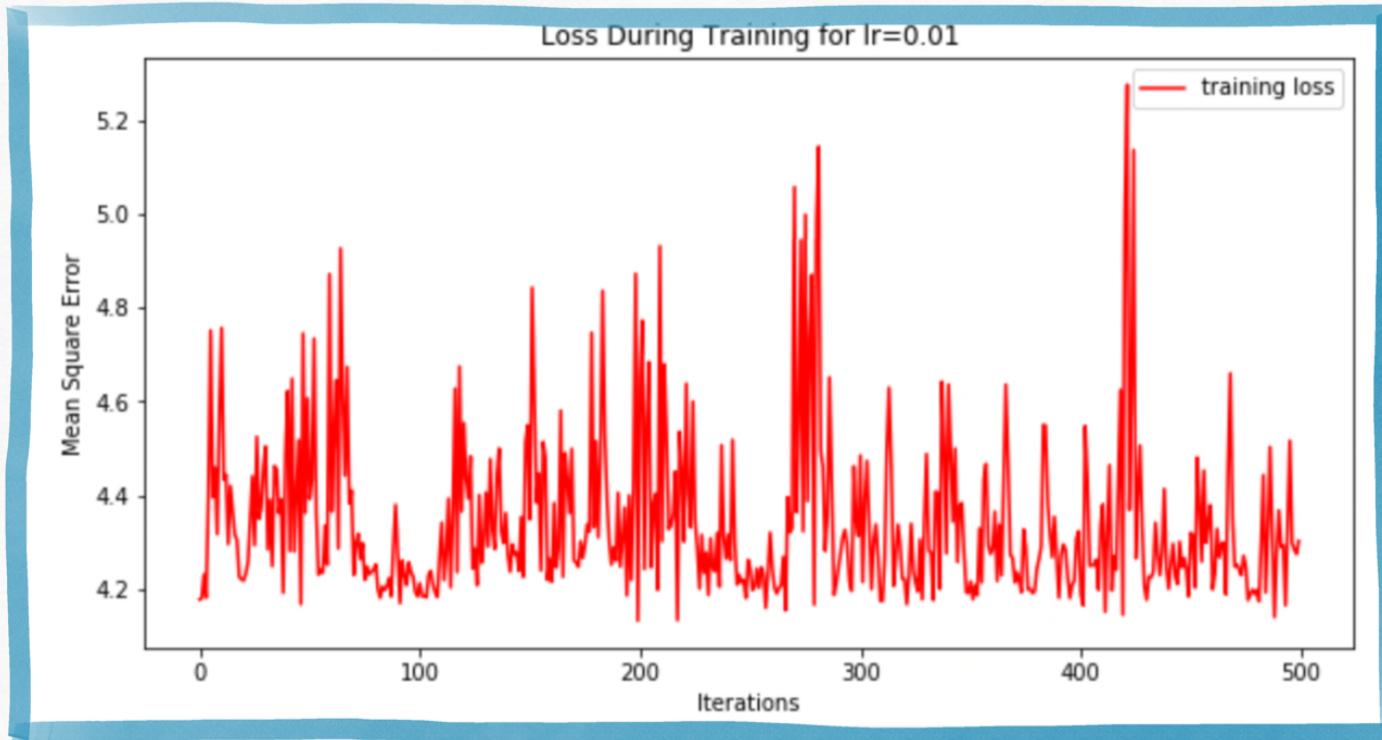
## DIAGNOSING ISSUES USING THE TRACE PLOT:

Can you tell by looking at the traceplot that training ended at a stationary point? Is our choice of the learning rate appropriate?



## DIAGNOSING ISSUES USING THE TRACE PLOT:

Can you tell by looking at the traceplot that training ended at a stationary point? Is our choice of the learning rate appropriate?



## DIAGNOSING ISSUES USING THE TRACE PLOT:

Can you tell by looking at the traceplot that training ended at a stationary point? Is our choice of the learning rate appropriate?

